


Models used to train and test the frozen PPO attacker



Each victim: 10,000 training images and 12 epochs | PPO: 1,000 policy images and 600 scheduled episodes

Final evaluation: 300 separate images and 25 calls per image

Seeds 17, 29, and 41 produce 9 separately trained PPO policies and 9 frozen evaluations